
Introduction to FastFlow programming

— Massimo Torquati —

<massimo.torquati@unipi.it>

SPM lecture 1

Master Degree in Computer Science
Master Degree in Computer Science & Networking
University of Pisa

Lessons calendar

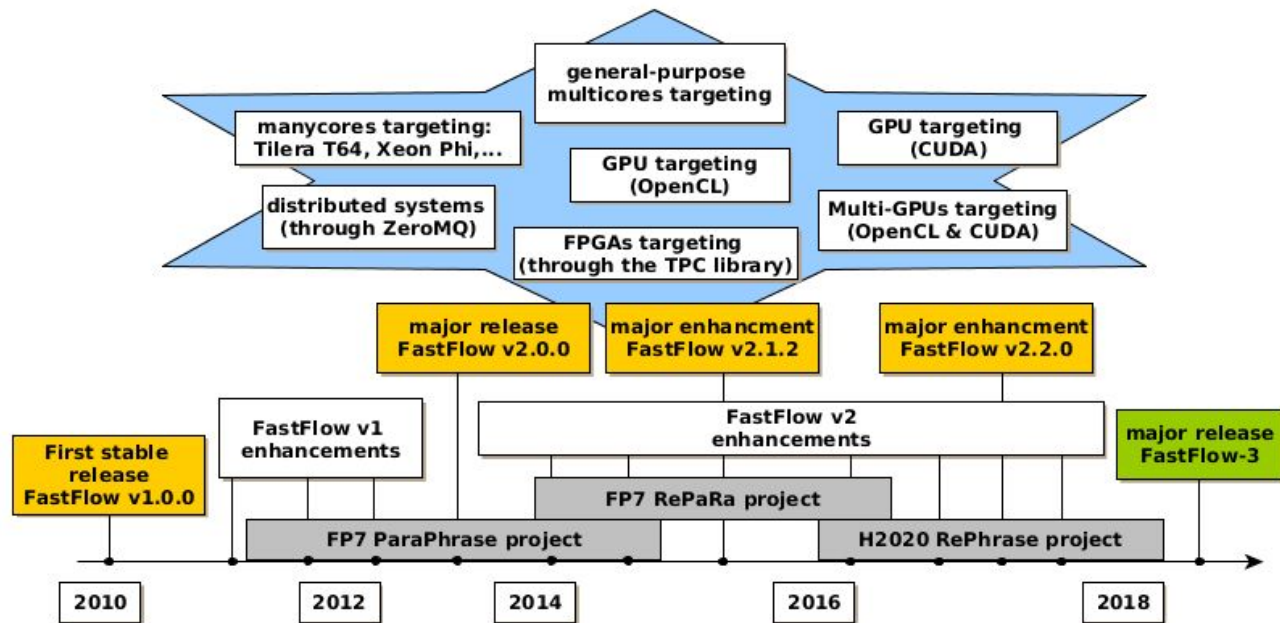
	When	What	Notes
Lesson1	Monday, April 12th	Introduction to FF, seq & pipeline	today!
Lesson2	Thursday, April 15th	Farm skeletons	
Lesson3	Monday, April 19th	Map, ParallelFor	
Lesson4	Thursday, April 22nd	Farm/ParallelFor examples, MDF, D&C patterns	
Lesson5	Monday, April 26th	FastFlow Building Blocks	9-10.30am (with a smaller break)
Lesson6	Monday, May 5th	WindFlow DSL example, Concurrency throttling in the ff_farm	
Lesson7	Monday, May 10th	Application examples (PARSEC, FastFilesCompressor)	
Lesson8	Monday, May 17th	Threads mapping, Memory allocators	To be confirmed!

What is FastFlow

- Research project started in 2010 (UniPI + UniTo)
- It is a parallel programming library written in C++ promoting structured parallel programming both at the RTS level as well as at the higher level
 - RTS level by using tiny components called ***building blocks***
 - Application level by using well-known ***parallel patterns***
- It aims to tackle the “3Ps” challenge for multi/many-core platforms:
Programmability, Performance, Portability
- Project actively maintained and utilized in several EU funded research initiatives
- Currently, FastFlow is moving to distributed systems
 - There exists an old proof-of-concept implementation of the FastFlow library for Clusters of Workstations based on ZeroMQ (not maintained anymore)
 - *Opportunities for master thesis on this topic!*

History

- About a decade of research
- 100+ research papers
- Several Ph.D. thesis
- Some IT industries tested/used FastFlow



Download

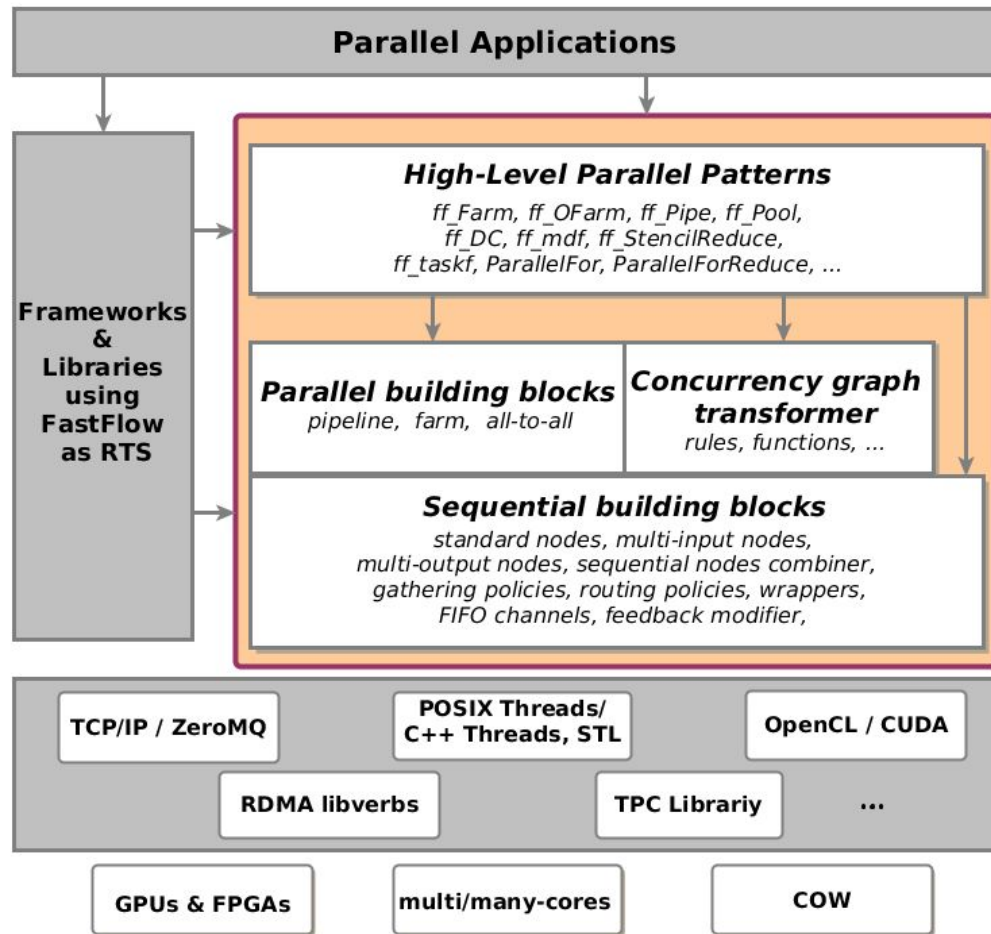
- The latest version (v3) available on GitHub (previously on Sourceforge):
 - git clone <https://github.com/fastflow/fastflow.git> it creates a fastflow directory
 - git pull inside the directory to get the latest update
 - **NOTE:** remember to run the Bash script *mapping_string.sh*
- It is possible to download the previous version (v2.2.0) containing simple examples and a tutorial (tab Releases on github)
- Notes:
 - From v2 to v3 the interface changed (minor modifications, mainly for the ff_farm)
 - To compile you need a recent C++ compiler
 - Compile with **-std=c++17** flag to avoid warnings
 - FastFlow is a class library (i.e., only header files) therefore do not forget to use the -O3 flag
 - You have to always include `<ff/ff.hpp>` first and then the others include needed (eg. `<ff/ParallelFor.hpp>`)

FastFlow stack

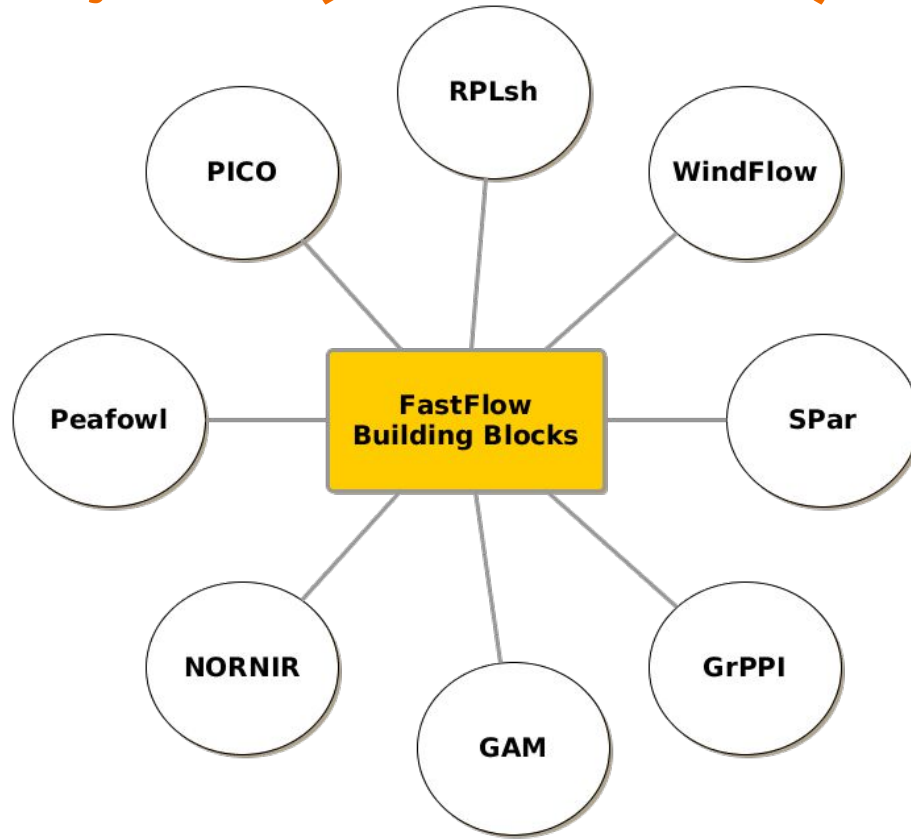
- Header-only library
- Multi-level APIs
- High-level Parallel Patterns targeting the application programmer
- Parallel and Sequential Building Blocks targeting the RTS programmer

Open-source LGPLv3

Home: <http://calvados.di.unipi.it/fastflow>
GitHub: <http://github.com/fastflow/fastflow>

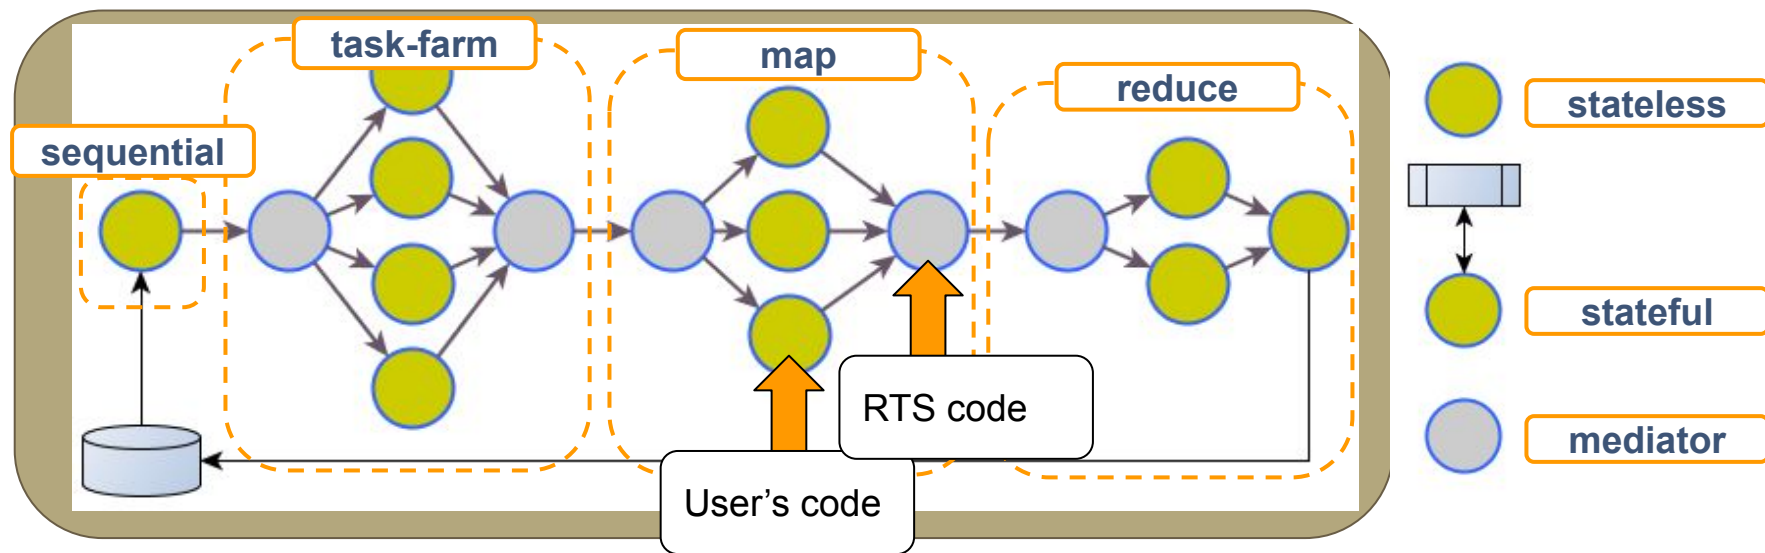


FastFlow ecosystem (research tools)



FastFlow programming model (basic concepts)

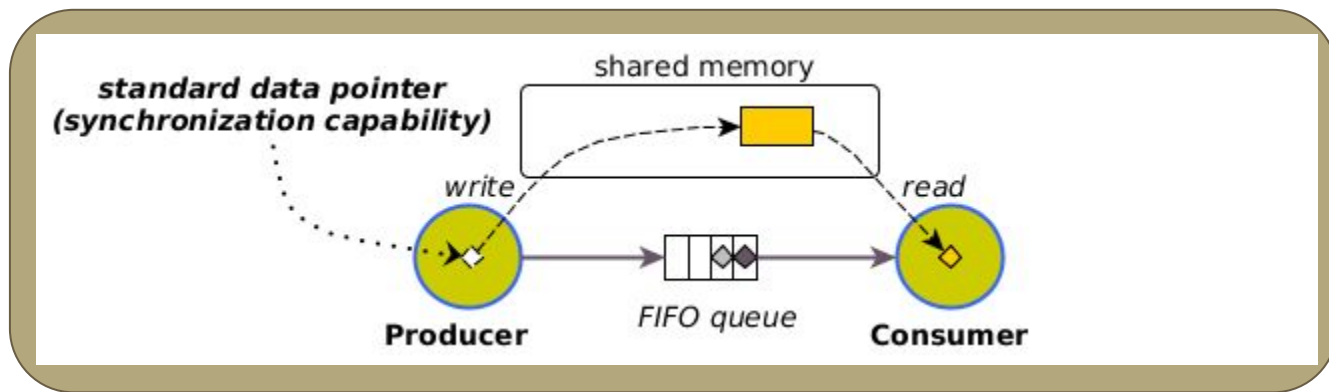
- An application is a **directed graph** whose **nodes** are computing entities (either sequential or parallel) and **edges** are channels carrying **data references (pointers)**



- Nodes can be **stateless** or **stateful**. In the most common case, a sequential node is a thread.

FastFlow programming model (basic concepts)

- Nodes synchronize through **message-passing** for accessing shared data
- Idea: *“Don’t communicate by sharing memory; share memory by communicating”*
 - This the Go language motto (<https://golang.org/doc/codewalk/sharemem/>)





- Writing a pointer into the FIFO queue enforces write ordering also on non TSO (Total Store Ordering) memory consistency models

FastFlow programming model (stream concept)

- **Data Stream** is a first class concept in the FastFlow library
- A stream is a sequence of data values (possibly infinite), coming from one or more sources. The items of a stream coming from the same source typically have the same data type.
- Streams can be either primitive ("*exo-stream*"), i.e. coming from a real source, for example a sensor, a network interface, a file, etc., or can be generated directly by the application ("*endo-streams*"), mainly by unrolling the iterations of a loop
- Terminology: stream *source/sink*, the first/last stage of a FastFlow pipeline

FastFlow programming model (stream concept)

- Streams coming from a real data source



- Data arrive with a given input rate
 - there could be fluctuations in the input rate, bursts, data losses, etc.
- The system should guarantee:
 - to sustain the maximum input rate
 - to process the single stream item within a given service time

FastFlow programming model (stream concept)

- Endo-stream: data stream produced by unrolling loops
 - The data source is a software module
 - Typically the length is bounded by an upper-bound
 - Typically generated at **full speed**, the most common objective is to reduce the overall completion time

```
for(i=start; i<stop; i+=step)
    allocate data for an item/task
    create a task
    send out the task to the next filter
```

FastFlow communication channels

- A communication channel connecting two FastFlow nodes is implemented by using a non-blocking SPSC FIFO queue (Single-Producer Single-Consumer First-In First-Out)
- The classical concurrency control approach for concurrent data structures is lock-based, which usually implies passive waiting when the queue is full or empty (**blocking** concurrency control protocol)
- SPSC queues are interesting because they can be implemented in a very efficient way on CMP platforms without the need of locks nor atomic instructions (i.e., CAS, LL/SC)
 - Lamport's circular buffer under Sequential Consistency memory model
- **Non-blocking** protocols use active-waiting loops with some back-off policies before trying again if the operation fails (e.g., push/pop operations in a concurrent queue)
 - **Obstruction-free progress**: a thread executed in isolation for a finite number of steps is capable to make progress and to terminate.
 - **Lock-free progress**: a failure/suspension of a thread does not prevent progress of other threads (there could be starvation, though). A lock-free algorithm is obstruction-free.
 - **Wait-free progress** (the strongest progress property): all concurrent operations in the algorithm are guaranteed to always complete in a finite number of steps. A wait-free algorithm is lock-free.

FastFlow communication channels

- FastFlow channels may have either **bounded** or **unbounded** capacity. They are implemented by the `uWSR_Ptr_Buffer` class (file `ff/ubuffer.hpp`) and we will refer to them as **uSPSC** queues
- The implementation of the uSPSC queue is based on the bounded non-blocking **SPSC** queue (file `ff/buffer.hpp` -- class `WSR_Ptr_Buffer`)
 - SPSC is basically a Lamport's buffer working also for Weak Memory Consistency Models
- FastFlow users do not have to deal with FastFlow channels!

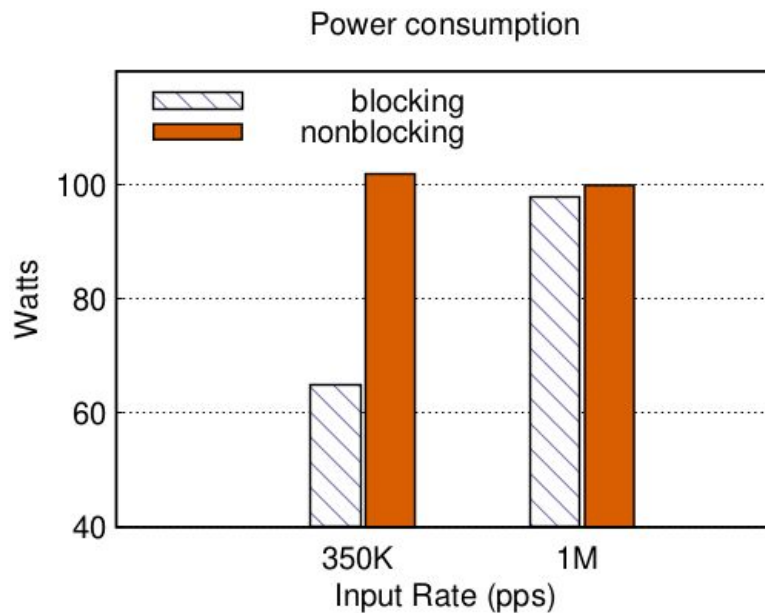
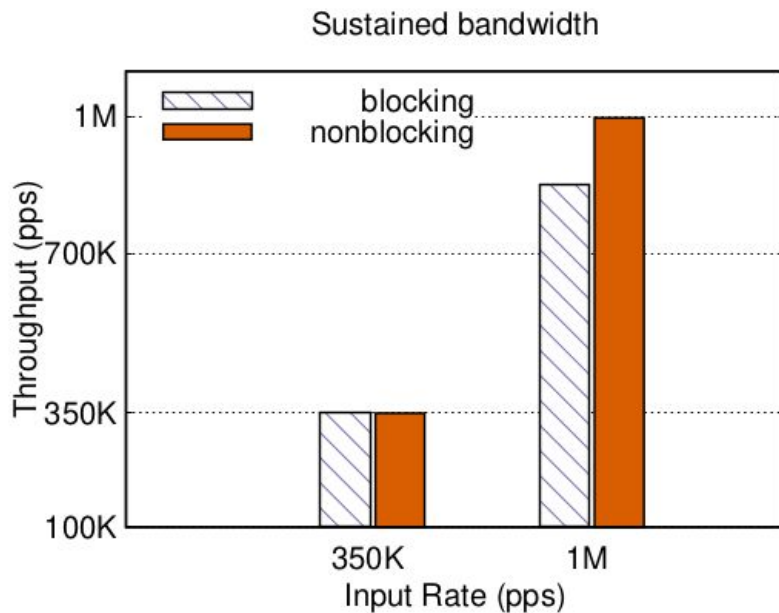
SPSC non-blocking FIFO buffer pseudo-code (uSPSC queues are a bit more complex)

```
1 bool push(void* data) {  
2     if (buf[pwrite]==NULL) {  
3         WMB(); // write-memory-barrier  
4         buf[pwrite] = data;  
5         pwrite+=(pwrite+1>=size)?(1-size):1;  
6         return true;  
7     }  
8     return false;  
9 }
```

```
10 bool pop(void** data) {  
11     if (buf[pread]==NULL)  
12         return false;  
13     *data = buf[pread];  
14     buf[pread]=NULL;  
15     pread+=(pread+1>=size)?(1-size):1;  
16     return true;  
17 }
```

FastFlow communication channels

- Performance vs Power Consumption (uSPSC)



- The blocking version protects push/pop operations with standard mutexes/CVs, and if the queue is empty the consumer will block

FastFlow communication channels




Xeon and KNL:

- mapping0: same core 2 contexts
- mapping1: 2 cores same CPU
- mapping2: 2 CPUs

Power:

- mapping0: same core 2 contexts
- mapping1 : 2 cores same CMP
- mapping2: same CPU 2 CMPs
- Mapping3: 2 CPUs

Details of the implementation of the uSPSC in:
M. Aldinucci, M. Danelutto, P. Kilpatrick, M. Meneghin, and M. Torquati, “An Efficient Unbounded Lock-Free Queue for Multi-core Systems” Euro-Par 2012, doi:10.1007/978-3-642-32820-6_65




High-Level Parallel Patterns

- **Stream parallel:** `ff_Pipe`, `ff_Farm`, `ff_OFarm`
- **Data parallel:** `ff_Map`, `ParallelFor*`, `poolEvolution`, `ff_stencilReuced*`
- **Task parallel:** `ff_DC`, `ff_mdf`

They try to address application
programmers' needs

FastFlow Name	Pattern description
<code>ff_Pipe</code>	Pipeline pattern modeling a data-flow <i>sequence</i> of sequential and parallel patterns. It allows to implements both linear and non-linear compositions of parallel patterns. The FastFlow implementation of this pattern is based on the <i>pipeline</i> building block.
<code>ff_Farm</code> , <code>ff_OFarm</code>	Task-farm pattern modeling functional replication. The <code>ff_OFarm</code> pattern guarantees to preserve input ordering. Both are built on top of the <i>farm</i> building block.
<code>ff_Map</code>	In this pattern a function F is applied to all elements of the input collection. The parallel calculation of the function F over disjoint partitions of the initial collection does not require any communication/synchronization except for a barrier at the end. It is implemented on top of the <i>farm</i> building block by using the broadcast and gather_all distribution and collection policies. The FastFlow library provides also an implementation based on the ParallelForReduce pattern.
<code>ff_mdf</code>	It models the Macro Data-Flow execution model. A program is interpreted by focusing on the functional dependencies among data [20, 69].
<code>poolEvolution</code>	This pattern models the evolution of a population according to the principles typical of evolutionary computing [12].
<code>ff_DC</code>	This pattern models <i>Divide & Conquer</i> computations [109].
<code>ParallelFor</code> , <code>ParallelForReduce</code>	They model data-parallel computations, allowing the parallelization of loops with independent iterations and also having reduction variables [119].
<code>ff_stencilReduceOCL</code> , <code>ff_stencilReduceCUDA</code>	These patterns model iterative stencil plus reduce computations targeting GPU accelerators either using OpenCL or using CUDA code [23, 8].

FastFlow sequential node (ff_node_t class)



```
virtual TOUT*  svc(TIN* task) = 0;    // encapsulates user's code
virtual int    svc_init();           // initialization code
virtual void   svc_end();             // finalization code
```

- **svc_init()** is called once at the beginning each time the thread associated to the node is (re-)started
- **svc()** is called for each input item present in one of the input queues. Its return value can be:
 - A memory pointer to the result data
 - **EOS** (End-Of-Stream)
 - **GO_ON**
 - A few other special values can be returned
- **svc_end()** is called before terminating the node (or when the node is put to sleep).
NOTE: It is not possible to send any “final” result into the output queues.

```
class myClass: public ff_node_t<IN_t, OUT_t>{
public:
    OUT_t* svc(IN_t * in) {
        <business logic code of the node
        producing the output task (out)>;
        return out;
    }
    int svc_init() { .....; return 0; }
    void svc_end() { ..... }

protected/private:
    <local-state, if any>
};
```

Sequential node life-cycle

- **svc_init()** is called once at the beginning each time the thread associated to the node is (re-)started
- **svc()** is called for each input item present in one of the input queues. Its return value can be:
 - A memory pointer to the result data
 - *EOS* (End-Of-Stream)
 - *GO_ON*
 - A few other special values can be returned
- **svc_end()** is called before terminating the node (or when the node is put to sleep).
NOTE: It is not possible to send any “final” result into the output queues.
- **eosnotify()** is called each time an EOS message is received by the node. It is possible to send data out into the output queue(s).

Sequential node life- cycle (simplified):

```
do {  
    if (svc_init() < 0) break;  
    do {  
        in = input_channel.pop();  
        if (in == EOS) {  
            // if this method has been redefined, the user's method  
            // is called and informed that the EOS has arrived  
            eosnotify();  
            output_channel.push(EOS);  
        } else {  
            out = svc(in); // it calls the business logic code  
            if (out == GO_ON) continue;  
            output_channel.push(out);  
        }  
    } while(out != EOS);  
    svc_end();  
} while(true);
```

Standalone node

- This is just an example to understand the life-cycle of a node
 - The FastFlow graph is composed by a single node
- The output is:
 - Hello. I'm going to start
 - Hi! (0)
 - Hi! (1)
 - ...
 - Hi! (5)
 - Goodbye!
- The method `run_and_wait_end` is needed to call the base class methods `run()` and `wait()` to start the thread associated to the node and to wait for the termination of the node, respectively.

```
#include <ff/ff.hpp>
using namespace ff;
struct myNode:ff_node_t<int> {
    int svc_init() {
        std::cout << "Hello. I'm going to start\n";
        counter = 0;
        return 0;
    }
    int* svc(int*) {
        if (++counter > 5) return EOS;
        std::cout << "Hi! (" << counter << ")\n";
        return GO_ON; // keep calling the svc method
    }
    void svc_end() { std::cout << "Goodbye!\n"; }
    // starts the node and waits for its termination
    int run_and_wait_end(bool=false) {
        if (run() < 0) return -1;
        return wait();
    }
    long counter;
};
int main() {
    myNode mynode;
    return mynode.run_and_wait_end();
}
```

Source and Sink nodes

Source node creating a stream of Tasks

```
struct Source: ff_node_t<Task> {  
    Task *svc(Task *) {  
        // generates N tasks and then EOS  
        for(long i=0;i<N; ++i) {  
            ff_send_out(new Task);  
        }  
        return EOS;  
    };  
};
```

Sink node absorbing the stream

```
struct Sink: ff_node_t<Task> {  
    Task *svc(Task * task) {  
        // do something with the task  
        do_Work(task);  
        delete task;  
        return GO_ON; // it does not send out task  
    };  
};
```


Pipeline

```
struct Source: ff_node_t<myTask> {  
    myTask *svc(myTask *) {  
        for(long i=0;i<10;++i)  
            ff_send_out(new myTask(i));  
        return EOS;  
    }  
};  
struct Stage: ff_node_t<myTask> {  
    myTask *svc(myTask *task) {  
        return task;  
    }  
};  
struct Sink: ff_node_t<myTask> {  
    myTask *svc(myTask* task) {  
        f3(task);  
        return GO_ON;  
    }  
};  
Source _1;  
Stage _2;  
Sink _3;  
ff_Pipe<> pipe(_1,_2,_3);  
pipe.run_and_wait_end();
```

- Pipeline stages are nodes (ff_node_t)
- A pipeline itself is a node
 - It is possible to build pipe of pipes
- A node, for each input, may generate zero, one, or many outputs
- The first stage is the *source* node
 - It generates 10 items and then the **EOS**
- The last stage is the *sink* node that executes function 'f3' and then **GO_ON**
- In this example, the middle stage is just a forwarder node

Pipeline example

- Computing the sum of the square of the first N numbers by using a pipeline



```
// 3-stage pipeline
ff_Pipe<> pipe( first, second, third );
pipe.run_and_wait_end();
```


```
// 1st stage
struct firstStage: ff_node_t<float> {
    firstStage(const size_t len):len(len) {}
    float* svc(float *) {
        for(long i=0;i<len;++i)
            ff_send_out(new float(i));
        return EOS; // End-Of-Stream
    }
    const size_t len;
};
```

Possible extention: think about how to avoid using many new/delete


```
// 2nd stage
struct secondStage: ff_node_t<float> {
    float* svc(float *task) {
        float &t = *task;
        t = t*t;
        return task;
    }
};
```

```
// 3rd stage
struct thirdStage: ff_node_t<float> {
    float* svc(float *task) {
        float &t = *task;
        sum +=t;
        delete task;
        return GO_ON;
    }
    void svc_end() { std::cout << "sum = " << sum << "\n"; }
    float sum = {0.0};
};
```



Multi-Input and Multi-Output nodes



- Single-Input Single-Output node
- `ff_node_t<IN_t, OUT_t>`
- `ff_send_out(data)`



- Single-Input Multiple-Output node
- **`ff_monode_t<IN_t, OUT_t>`**
- `ff_send_out_to(data, channel-id)` channel-id goes from 0 to N-1



- Multiple-Input Single-Output node
- **`ff_minode_t<IN_t, OUT_t>`**
- `ff_get_channel_id()` to know from which channel we received the input (from 0 to N-1)
- *`fromInput()`* return true if the input comes from one of the input channels (useful in case of inputs coming both from input and feedback channels)