
Introduction to FastFlow programming

— Massimo Torquati —

<massimo.torquati@unipi.it>

SPM lecture 3


Master Degree in Computer Science
Master Degree in Computer Science & Networking
University of Pisa

Data Parallel (DP) Computations (recap)

- Large collections of data are partitioned among the number of computing resources (threads, processes, nodes, etc.) each one executing the same function F over the assigned partition of the collection
- The computation may be inplace (typical in *map-based* computations)
- Usually the so-called “owner computes rule” is applied (typical in *stencil-based* computations)
 - On shared-memory systems, the function F has read/write access to the elements of the assigned partition and read-only access to (some) other elements of the input collection
- The main goal of data parallel computations is to reduce the *completion time* to compute the entire collection
- Usually they are encountered in sequential programs as loop-based computations

Data Parallel (DP) Computations (recap)

- Logical implementation of a DP computation




Data Parallel Computations in FastFlow

- Data parallel patterns are: Map, Reduce, Map+Reduce, Stencil, Stencil+Reduce
- In FastFlow all DP patterns have been implemented by using the ParallelFor/ParallelForReduce patterns
 - The Map pattern in FastFlow is based on the ParallelForReduce. It implements a pure Map and a Map+Reduce
 - We will see that the Map pattern can be used as a Worker of a farm and/or as a stage of a pipeline.
 - We will not consider the StencilReduce pattern

ParallelFor

- The *ParallelFor* pattern can be used to parallelize loops with independent iterations
- The class interface is defined in the file *ff/parallel_for.hpp*
- Example:

```
// A and B are 2 arrays of size N  
  
for(long i=0; i<N; ++i)  
    A[i] = A[i] + B[i];
```



```
#include <ff/parallel_for.hpp>  
using namespace ff;  
  
ParallelFor pf(8); // defining the object  
  
pf.parallel_for(0, N, 1, 0, [&A,B](const long i) {  
    A[i] = A[i] + B[i];  
}, 4);
```

- Constructor interface (all arguments have a default value)
 - *ParallelFor(maxnworkers, spinWait, spinBarrier);*
- On the basis of the number and types of arguments you have different *parallel_for* signatures
 - *parallel_for(first-index, last-index, stepsize, chunksize, bodyFunction, nworkers);*
 - The *bodyFunction* is a C++ lambda

Example: sum of the square

- Previously implemented generating a stream of N floats and using a pipeline+farm schema
- Basically we unfolded the loop:

```
sum=0.0;  
for(float i=0.0; i<N; i++) sum += i*i;
```

- Let's use now a ParallelFor for computing the square of the elements of a `std::vector<float>`
- The reduction (i.e. the total sum) will be computed sequentially
 - See the `pf_square.cpp` file

ParallelForReduce

- The *ParallelForReduce* pattern can be used to parallelize loops with independent iterations and having reduction variables (map+reduce pattern)

- Example:

```
// A is an array of long integers of size N
long sum = 0;
for(long i=0; i<N; ++i)
    sum += A[i];
```



```
#include <ff/parallel_for.hpp>
using namespace ff;

ParallelForReduce<long> pfr;
long sum=0;
pfr.parallel_reduce(sum, 0,
                    0,N,1,0, [](const long i, long &mysum) {
                        mysum += A[i] + B[i];
                    },
                    [](long &s, const long e) { s += e; }
);
```

- Constructor interface is the same of the *ParallelFor* (but the template type)
- *parallel_reduce* method interface:
 - *parallel_reduce*(var, identity-val, first, last, step, chunksize, mapF, reduceF, nworkers);
 - *mapF* and *reduceF* are C++ lambdas

Example: sum of the square

- Let's use a ParallelForReduce for computing the square and the reduction
 - No array needed we can parallelize directly the initial loop

```
sum=0.0;  
for(float i=0.0; i<N; i++) sum += i*i;
```

- See pf_square2.cpp

Example: dot product


- The dot product operation (or scalar product, inner product) takes two vectors (A, B) of the same length, it produces in output a single element that is the sum of the products of the corresponding elements of the two vectors. Example:

```
std::vector<double> A(N), B(N);  
..... // A and B initialized  
double s=0.0;  
for(size_t i=0; i<N; i++) s += A[i]*B[i];
```

- Let's have a look at the test_dotprod_parfor.cpp inside the tests directory


ParallelFor* implementation

- The ParallelFor* patterns are implemented on top of the farm building block
- The skeleton is a master-worker, where the master is the Scheduler of loop iterations
- The Scheduler can be disabled by calling the ParallelFor* method *disableScheduler()*



Iterations scheduling

- Let's consider the following case
 - `for(size_t i=0;i<N;++i) A[i] = F(A[i]); //map-like computation`
- The time difference for computing distinct elements in the array A may be large (or very large)
- Problem: How do we schedule the loop's iterations?




Iterations scheduling

Suppose to have 3 workers and a chunksize=2, then the initial plan used for scheduling iterations is

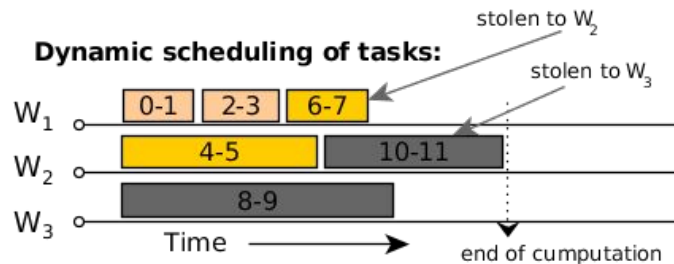
wid	#tasks	min-max
0	2	0-3
1	2	4-7
2	2	8-11



Static assignment of tasks:



Dynamic scheduling of tasks:




Iterations scheduling in the ParallelFor* patterns

- Iterations are scheduler according to the value of the *chucksize* parameter
 - *parallel_for(first-index, last-index, stepsize, **chucksize**, bodyFunction, nworkers);*
 - *parallel_reduce(var, identity-val, first, last, step, **chucksize**, mapF, reduceF, nworkers);*
- Three options:
 1. **chucksize = 0** : static scheduling
 - Each worker gets a contiguous chunk of $\sim(\text{\#iteration} / \text{\#workers})$ iterations
 2. **chucksize > 0** : dynamic scheduling with task granularity equal to *chucksize*
 3. **chucksize < 0** : static scheduling with task granularity equal to *chucksize*, chunks are assigned to workers in a round-robin fashion


Mandelbrot set example

- Simple DP computation:
 - Each pixel of the image can be computed independently
 - Black pixels require much more iterations
 - Easy to implement by using a farm or a ParallelFor
- A static partitioning of the image quickly leads to unbalanced computation among Workers
- Example:
 - Image size 2048x2048, max iterations per point 10^3 , 48 Workers, 24 cores (2-way HT)
 - Static partitioning of rows (i.e. chunksize=0) **Max Speedup ~14**
 - Dynamic partitioning of rows (i.e. chunksize=1) **Max Speedup ~37**
- Code: mandel.cpp



Combining DP and Stream Parallel computations

- If one of the stages of a pipeline is a bottleneck its service time can be reduced by parallelizing the stage
- If the stage computes a Map or a Map+Reduce it can be parallelized by using a ParallelFor* pattern



- Two options:
 - Using a ParallelFor* pattern inside the svc method
 - Replacing the node with a ff_Map pattern

The FastFlow Map Pattern

- The FastFlow Map Pattern (`ff_Map`) is just a single-input single-output node that wraps a `ParallelForReduce` pattern
 - `ff_Map<IN_t, OUT_t, reduce-variable-type>`
- Inside a pipeline or a farm, it is generally better to use the `ff_Map` than a plain `ParallelForReduce` because some optimizations are automatically introduced by the Map (mapping of worker threads, scheduler disabled, etc..)

```
#include <ff/map.hpp>
using namespace ff;

struct myMap: ff_Map<Task,Task,float> {
    using map = ff_Map<Task,Task,float>;

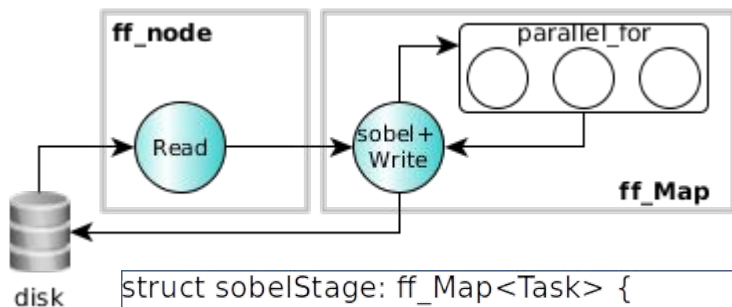
    Task *svc(Task *input) {

        map::parallel_for(...);

        float sum = 0;
        map::parallel_reduce(sum, 0.0, ....);

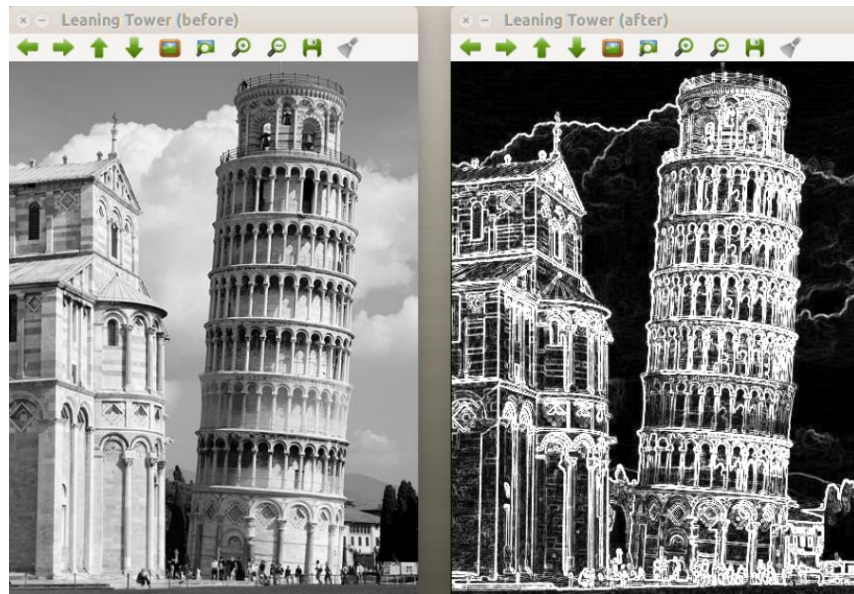
        return out;
    }
};
```


Example: Image filtering



```
struct sobelStage: ff_Map<Task> {
  sobelStage(int mapwrks):
    ff_Map<Task>(mapwrks, true) {};


  Task *svc(Task*task) {
    Mat src = *task->src, dst= *task->dst;
    ff_Map<>::parallel_for(1,src,src.row-1,
      [src,&dst](const long y) {
        for(long x=1;x<src.cols-1;++x) {
          .....
          dst.at<x,y> = sum;
        }
      });
    const std::string outfile="./out"+task->name;
    imwrite(outfile, dst);
  }
}
```



- The first stage reads some image from the disk
- Each image is converted in B&W and sent to the next stage
- The second stage applies the **Sobel filter** to each image by using OpenCV and then writes the result into the disk

Example: Image filtering

- We can parallelize the second stage by using a farm because the image are independent
- We can also normalize the farm (see the figure)
 - The emitter schedules file names of the image using an on-demand policy,
 - the Worker reads the image in parallel, then applies the filter using a Map and finally writes the resulting image on the disk (with a different name)
- There are two levels of parallelism: th n. of farms Workers, and the n. of Map Workers, how to balance them?



Example: Sobel filter

- Machine: 2 CPUs Xeon E5-2695 @2.4GHz eche CPU has 12 cores 2-way HyperThreading
- 320 images of different size (from a few KBs to some MBs)
- Results:**
 - a) sequential ~1m;
 - b) pipe(seq, map(4)) ~15s;
 - c) farm(map(4), 8) ~5s;
 - d) farm(seq, 32) ~3s
- The d) version is the *normal form*, it provides always the best performance (if it can be used)

